



Distributed Logs


Michael Enudi

Journey through the world of databases and data engineering




Scope

- Data Architectures
- Introduction to Distributed Logs
- Use case: Apache Kafka
- Apache Zookeeper
- Stream processing
- Stream Processing II (Kafka Streams)
- Stream analytics
- Lab: NYC Taxi and Fares
- Streaming: To be or not to be




**In software or data engineering, we
encounter data in two general states**



Data-at-rest

Batch

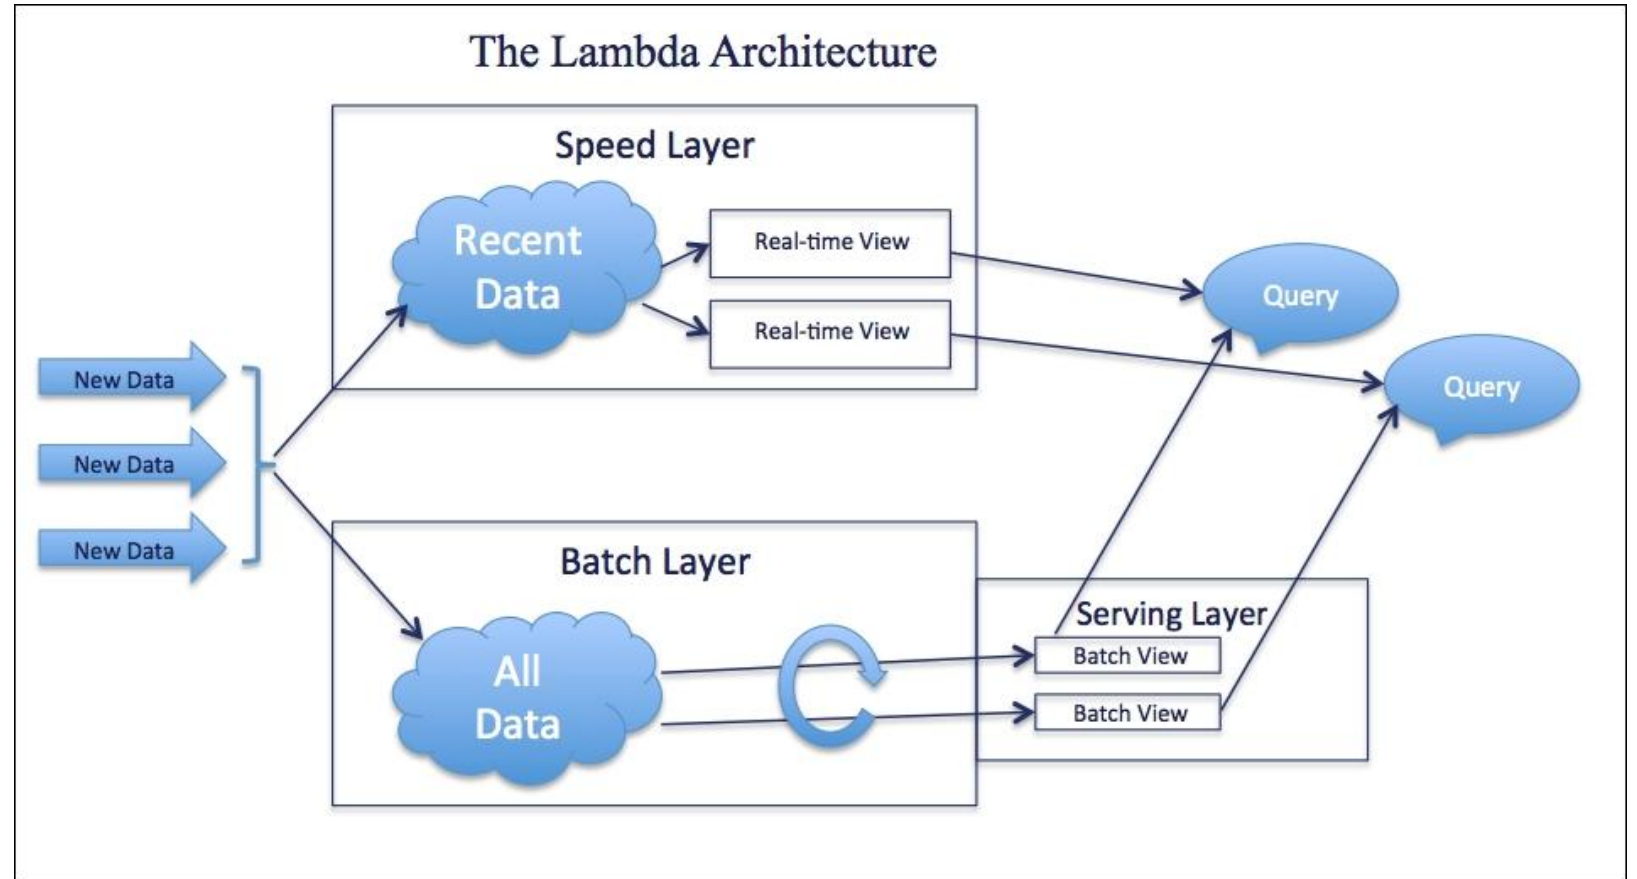


Data-in-motion


Stream/Streaming

“Data-in-Use” / “Serving”

Data architecture (Lambda)



Data architecture (Kappa)



Message Brokers



Distributed Logs





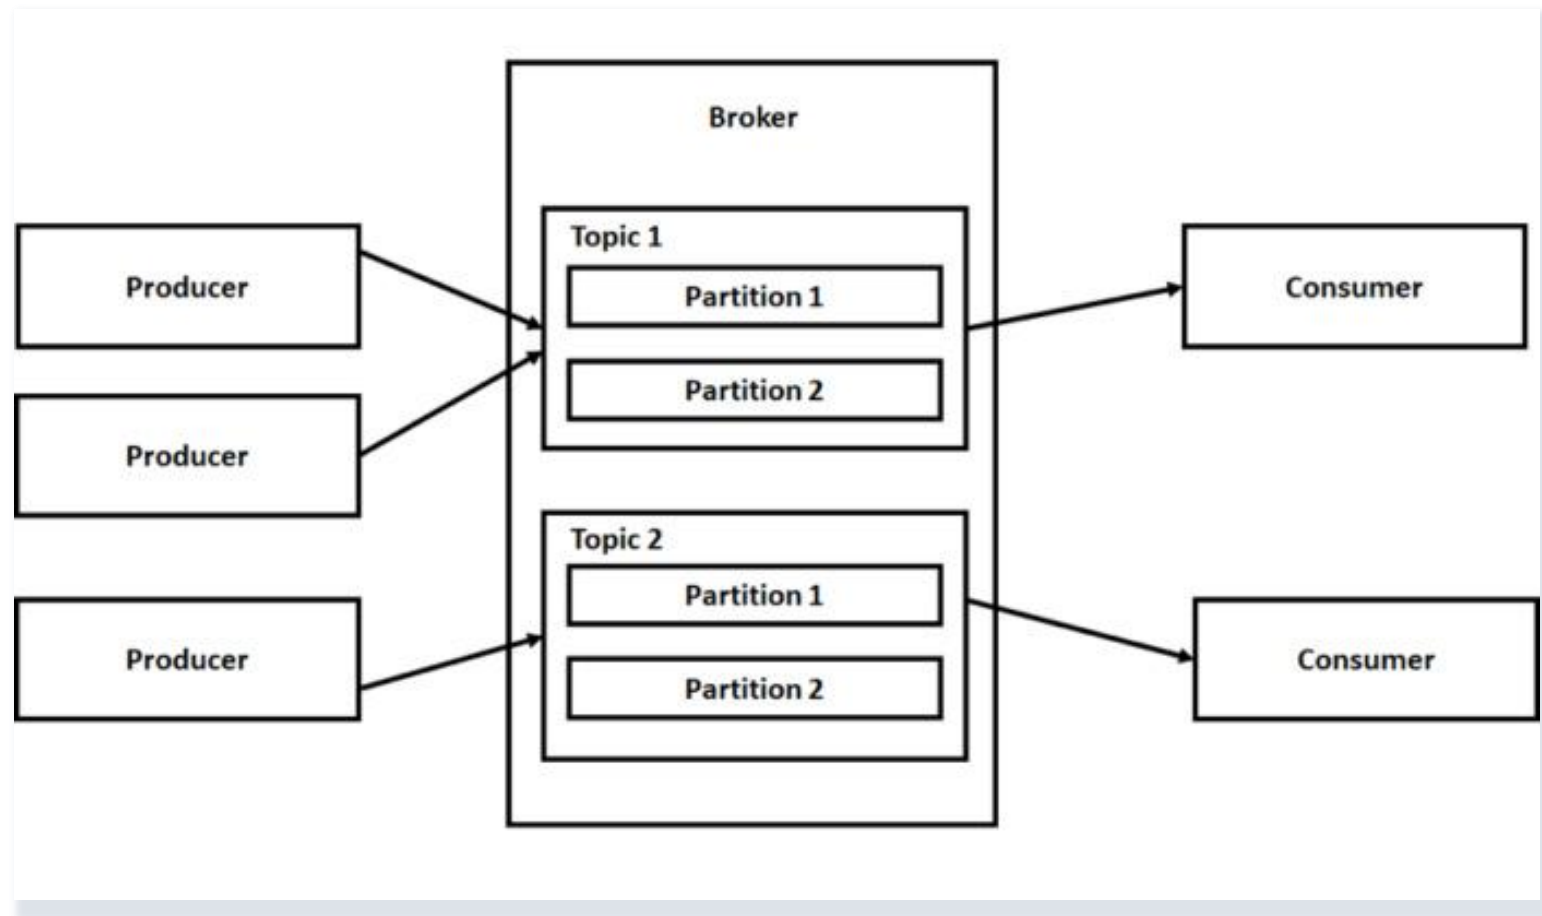
DISTRIBUTED COMMIT LOG AS A STORAGE PLATFORM



- ☐ Used as a Message Broker or as a building real-time data pipelines for streaming applications
- ☐ Kafka started from within LinkedIn in January 2011
- ☐ Current version as at June 2019 is 2.3.0
- ☐ Written in Scala and Java
- ☐ Uses Zookeeper (just like HBase) for distributed coordination and metadata repository.
- ☐ Confluent currently offers Kafka as part of the Confluent Platform.
- ☐ Use cases include message broker, streaming application, event sourcing, commit logging, activity tracking and more.
- ☐ High profile users of kafka include LinkedIn, Apple, eBay, Netflix, Yelp, Oracle, Uber, Citibank, Spotify and a host of others





- Topic
- Records / Messages
- Producers
- Consumers



Kafka Topic

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
---	---	---	---	---	---	---	---	---	---	----	----	----	----	----



Topic Partition Replication

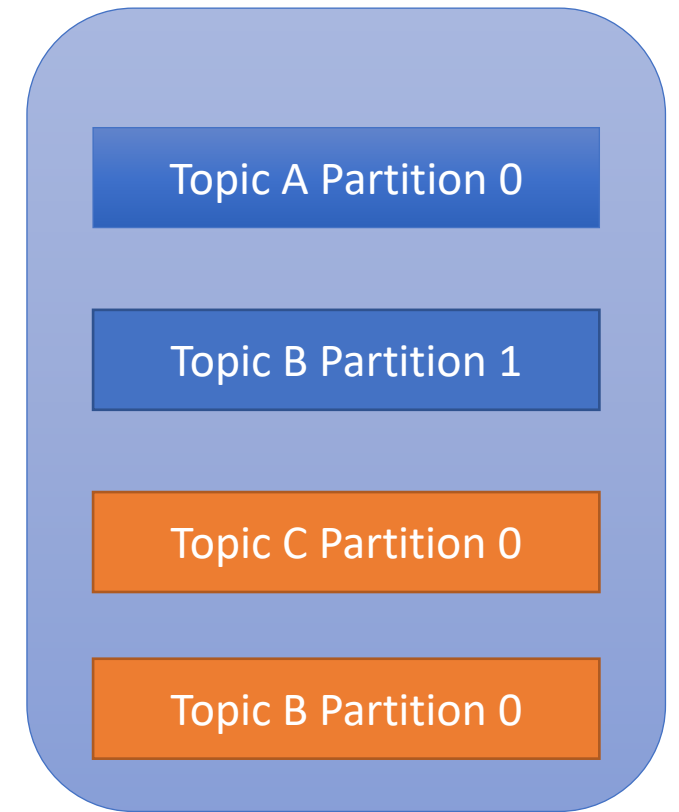
Topic A - 1 Partition & Replication factor of 3
Topic B - 2 Partition & Replication factor of 3
Topic C - 1 Partition & Replication factor of 2
Topic D - 1 Partition & Replication factor of 1



Broker 1



Broker 2



Broker 3

Kafka Messages

Topic

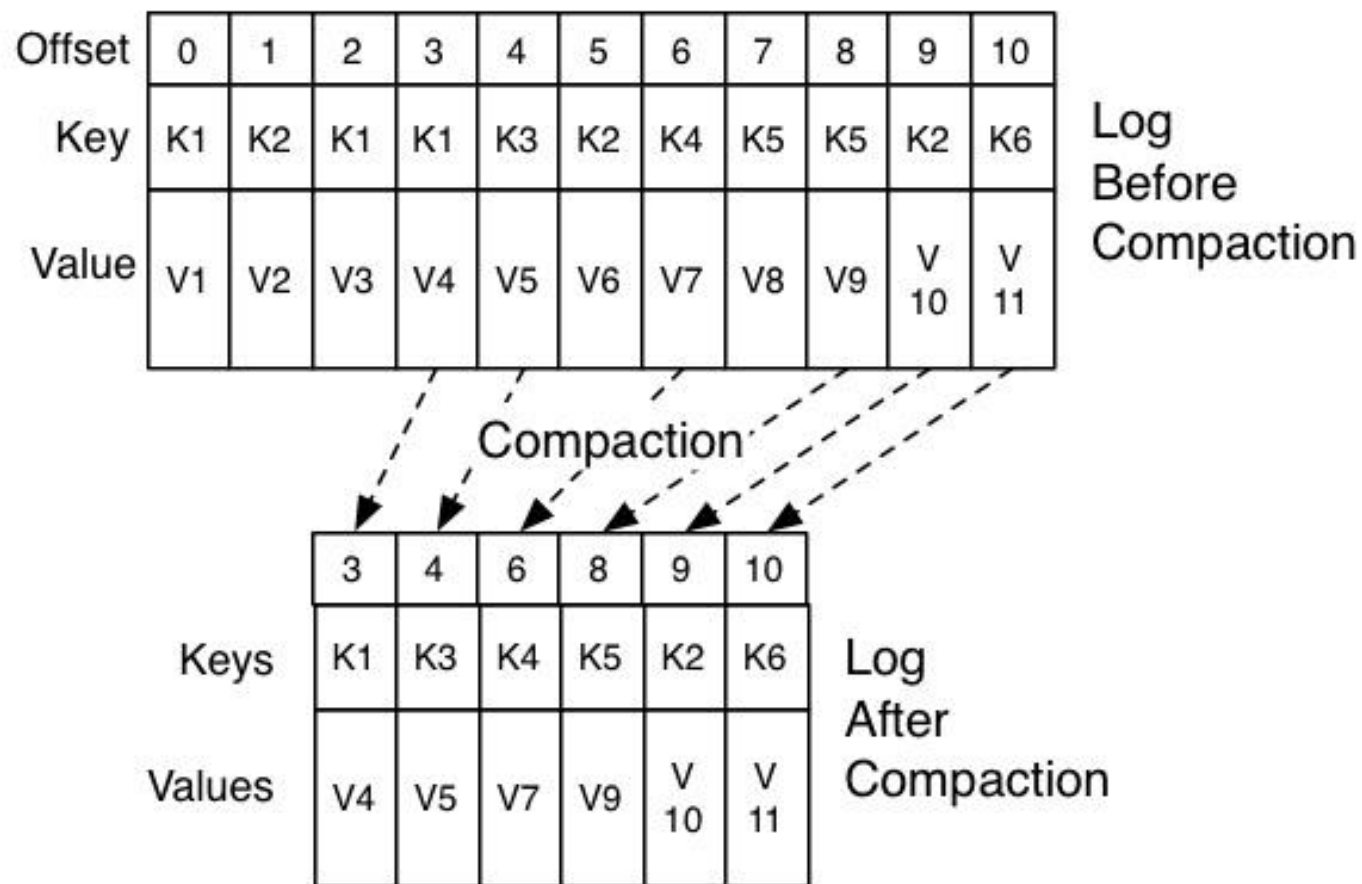
0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
---	---	---	---	---	---	---	---	---	---	----	----	----	----	----

Relational Table


0
1
2
3
4
5
6
4
7
8
9
10
11
12
13
14

Kafka Messages

Compaction Clean up policy

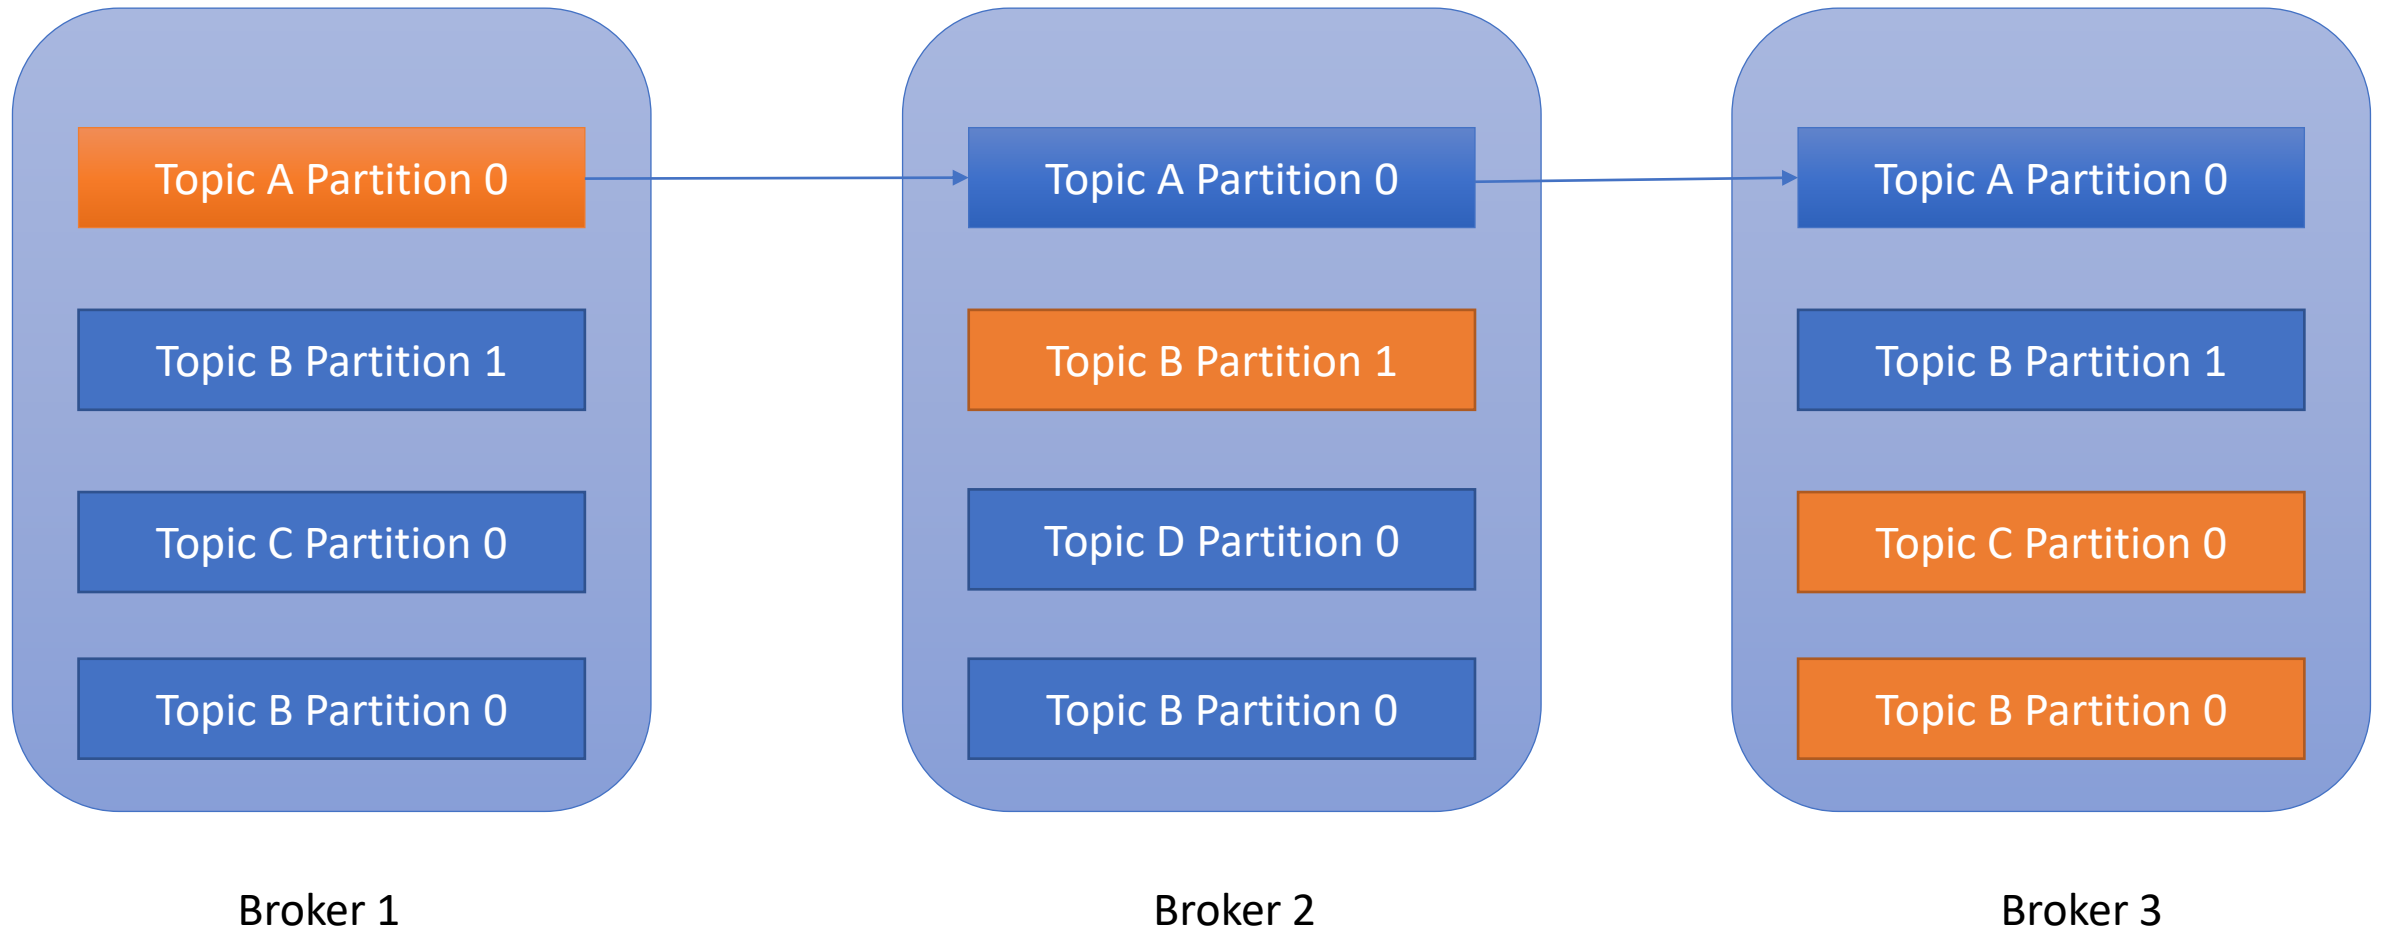


Kafka Producer

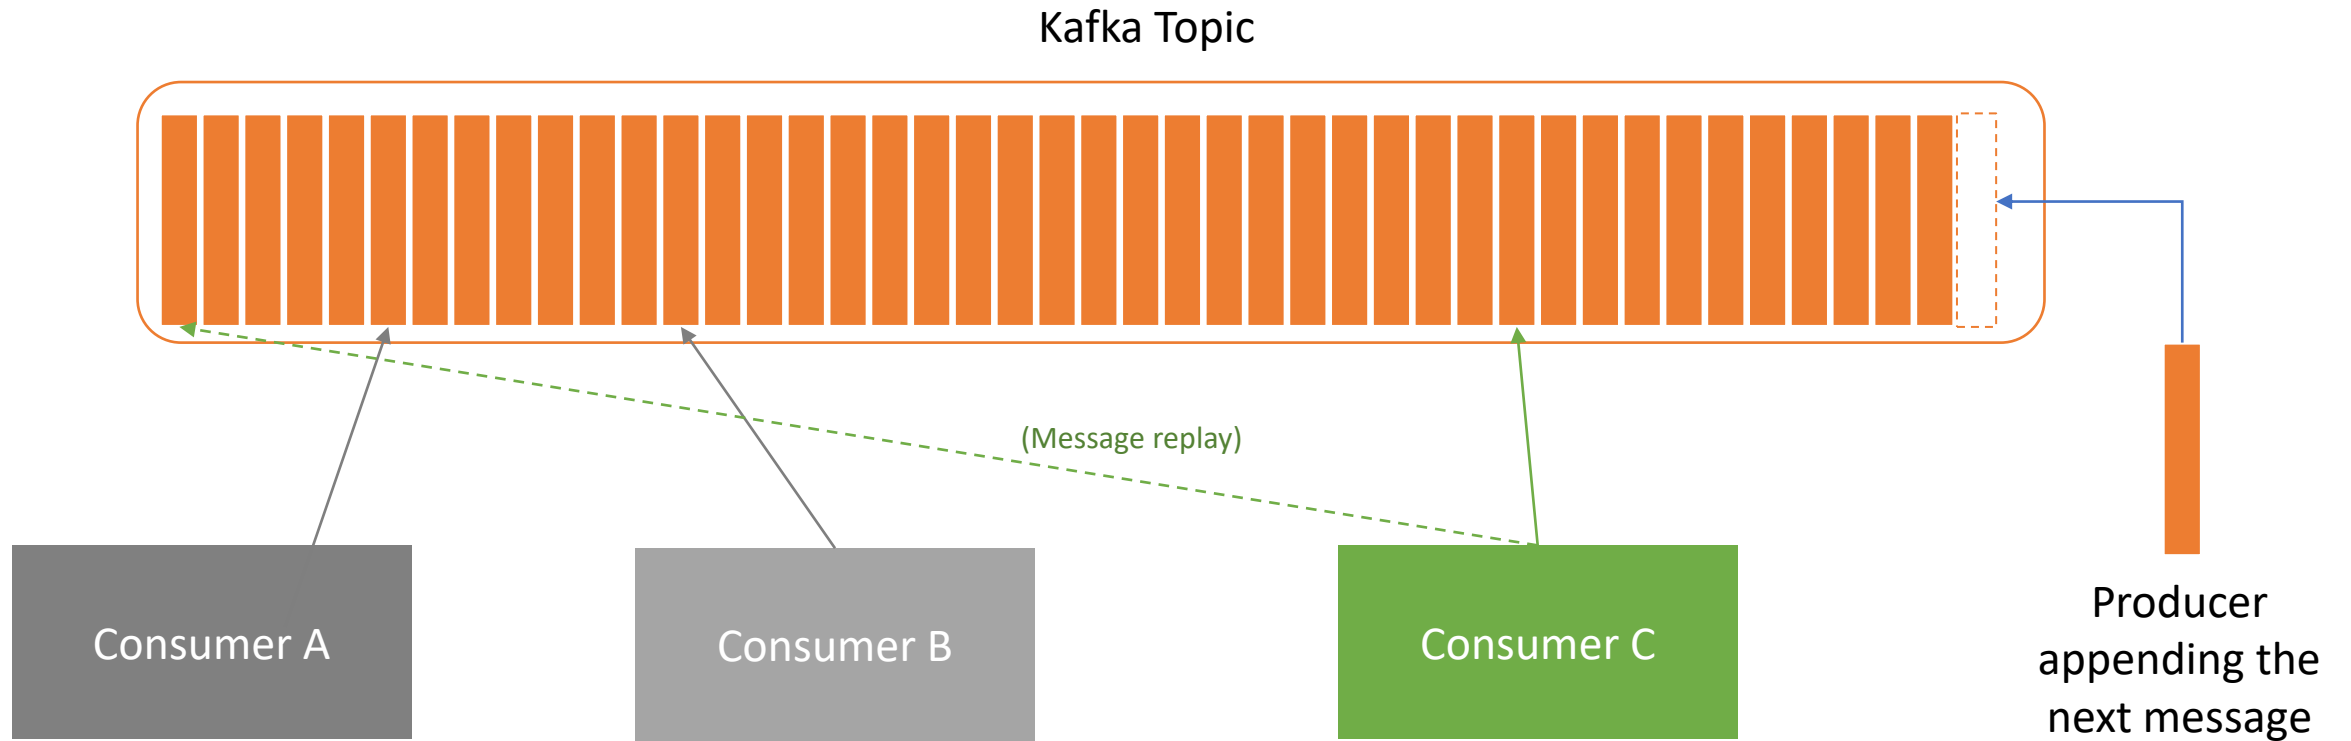


Kafka Producer Acknowledgement

Topic A - 1 Partition & Replication factor of 3
Topic B - 2 Partition & Replication factor of 3
Topic C - 1 Partition & Replication factor of 2
Topic D - 1 Partition & Replication factor of 1

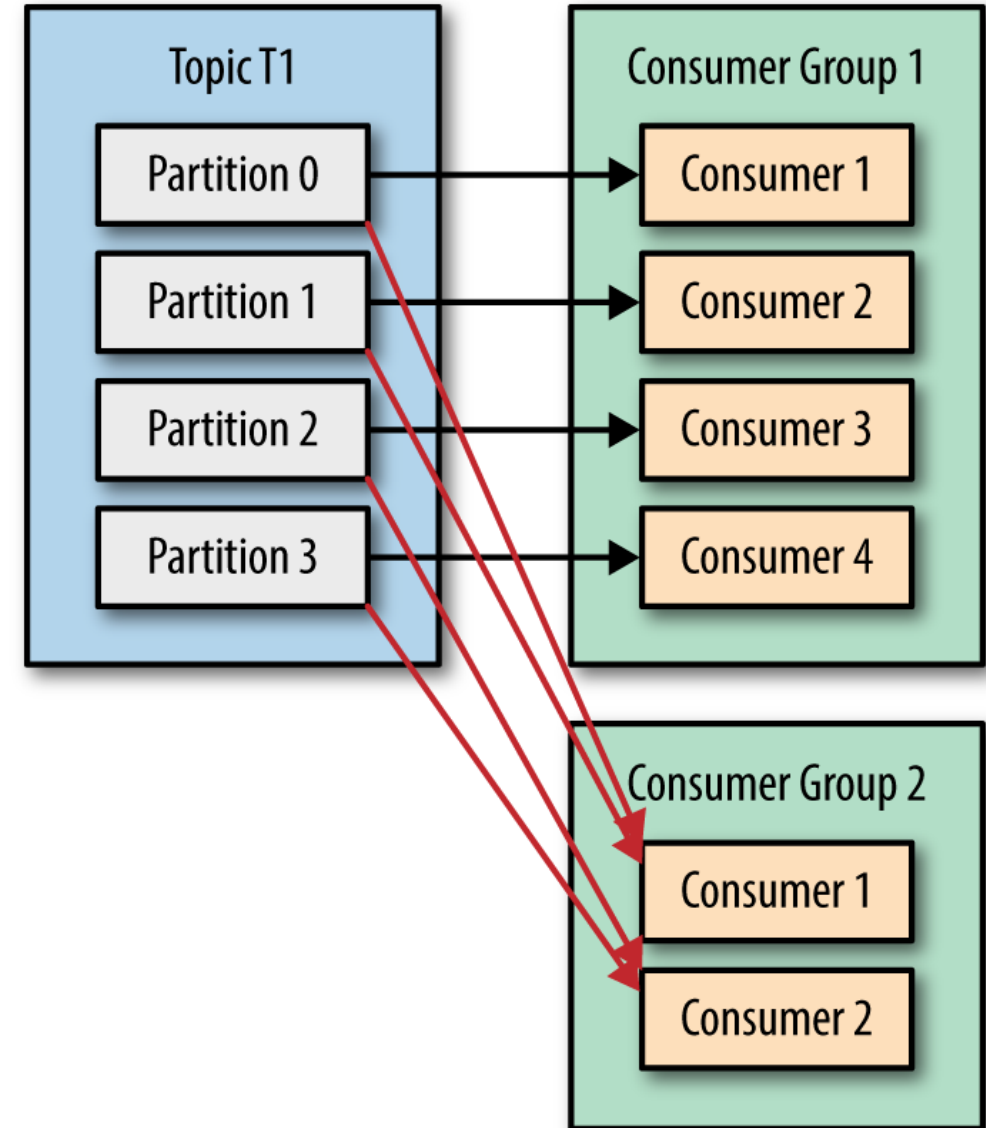



Kafka Consumers



Consumers subscribe to Topics and read from them at their various offsets

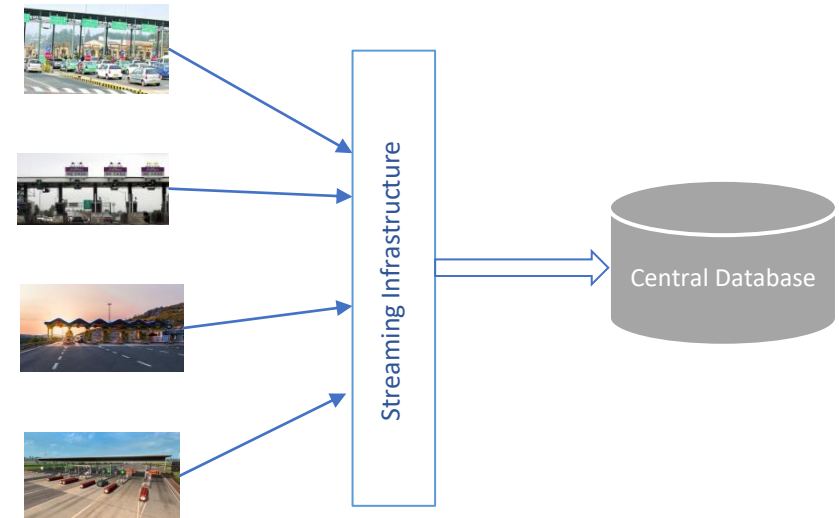
Kafka Consumer Groups






Use case

Toll payment collection and monitoring platform



KAFKA CONNECT + KAFKA STREAMS

KAFKA CONNECT + STREAMS

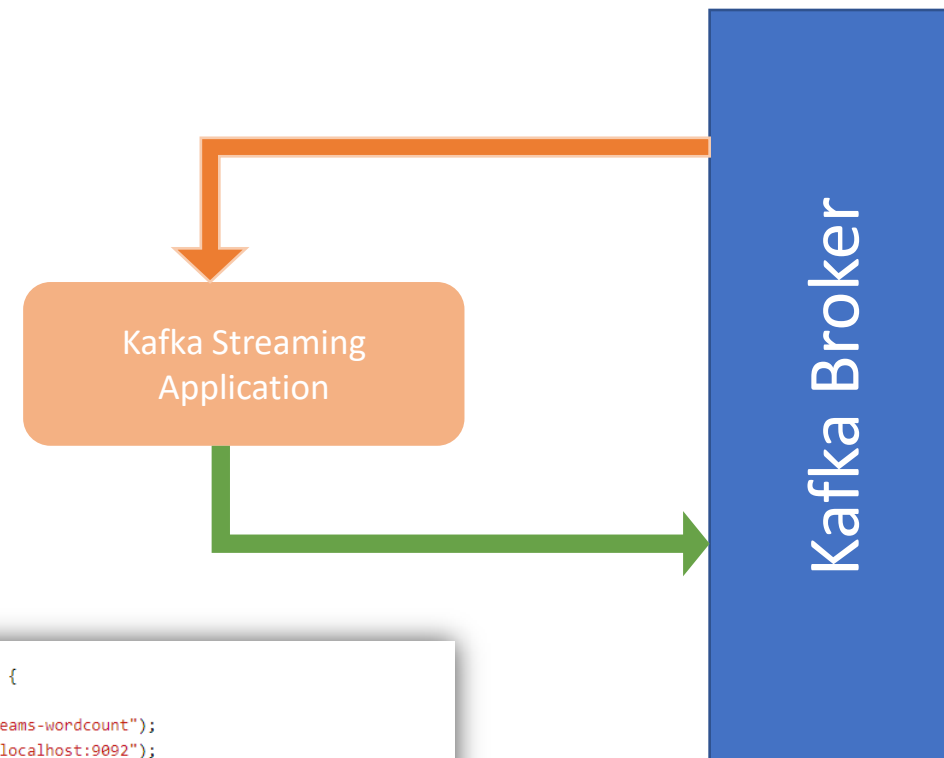


Kafka Stream

A library for implementing scalable stream processing applications that react to events in real-time in Scala and Java.

It is more than just executing a function between a consumer and a producer.

- ❑ Stream abstraction
- ❑ Streams-Tables duality
- ❑ Time
 - ❑ Windowing
- ❑ Aggregations
- ❑ State



```
public static void main(String[] args) throws Exception {
    Properties props = new Properties();
    props.put(StreamsConfig.APPLICATION_ID_CONFIG, "streams-wordcount");
    props.put(StreamsConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
    props.put(StreamsConfig.DEFAULT_KEY_SERDE_CLASS_CONFIG, Serdes.String().getClass());
    props.put(StreamsConfig.DEFAULT_VALUE_SERDE_CLASS_CONFIG, Serdes.String().getClass());

    final StreamsBuilder builder = new StreamsBuilder();

    KStream<String, String> source = builder.stream("streams-plaintext-input");
    source.flatMapValues(value -> Arrays.asList(value.toLowerCase(Locale.getDefault()).split("\\W+")))
        .groupByKey((key, value) -> value)
        .count(Materialized.<String, Long, KeyValueStore<Bytes, byte[]>>as("counts-store"))
        .toStream()
        .to("streams-wordcount-output", Produced.with(Serdes.String(), Serdes.Long()));

    final Topology topology = builder.build();
    final KafkaStreams streams = new KafkaStreams(topology, props);
    final CountDownLatch latch = new CountDownLatch(1);


    // ... same as Pipe.java above
}
```

A library for reusable component from export data from a Kafka topic to an external destination or for import data from an external source into Kafka.

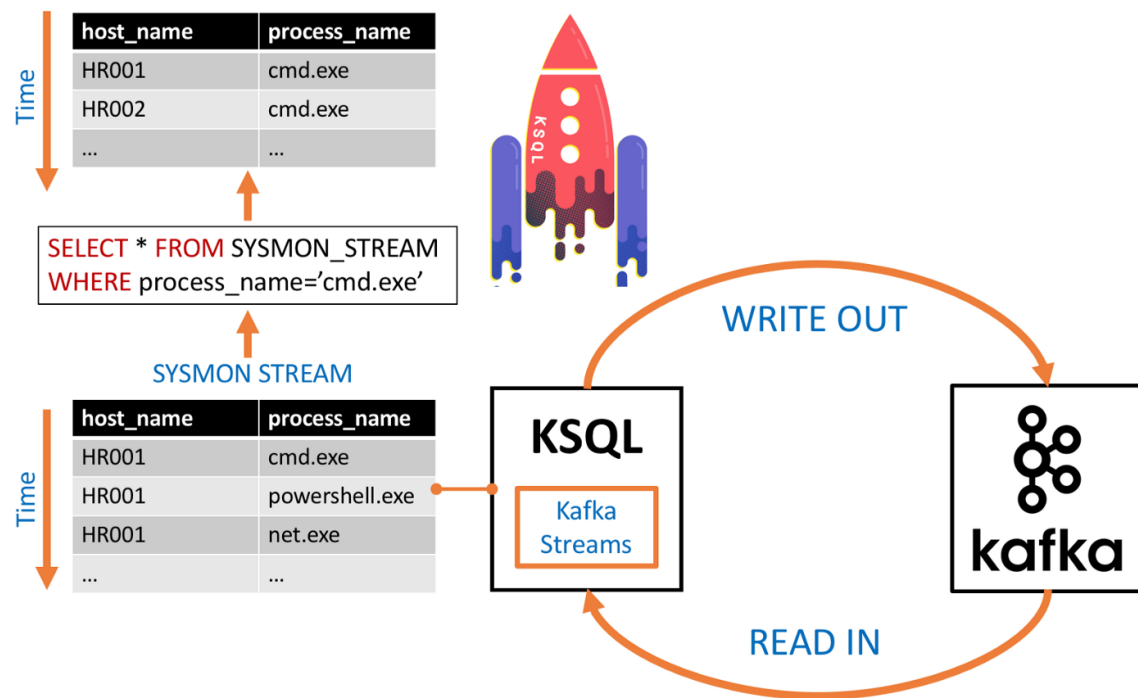
The name for this reusable component is called a connector.



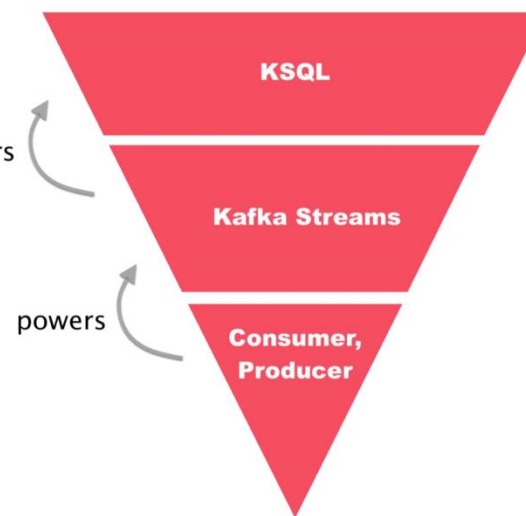
KAFKA CONNECT



KSQL



powers



ease of use

```
CREATE STREAM, CREATE TABLE,  
SELECT, JOIN, GROUP BY, SUM, ...
```

```
KStream, KTable,  
filter(), map(), flatMap(),  
join(), aggregate(), ...
```

```
subscribe(), poll(), send(),  
flush(), beginTransaction(), ...
```

flexibility

A thin vertical line is positioned to the left of the text.

DISTRIBUTED LOG WRAP-UP