


Document-oriented Databases


Michael Enudi

Journey through the world of databases and data engineering



Scope


- Introduction to Document Databases
- Use Case: MongoDB
- Movielens in MongoDB
- Data processing in MongoDB
- Data Analysis in MongoDB
- Lab: Anonymized MOT tests and results
- Document-Oriented Databases: To be or not to be.



MongoDB

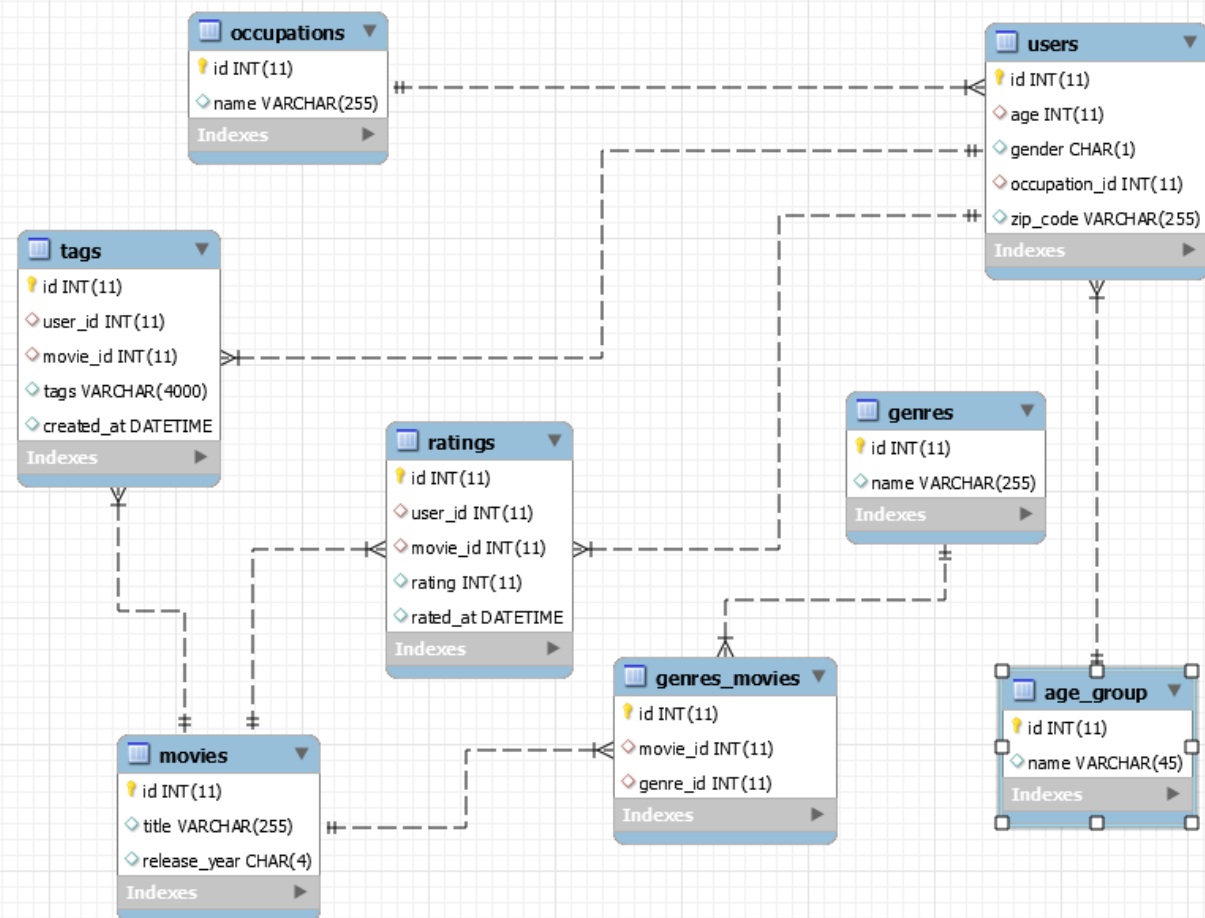
- Free and open source, linearly scalable Document-oriented database
- Most popular NoSQL Database as of May, 2019.
- Most popular Document-oriented database
- Written in C++
- Actively supported by MongoDB Inc.
- First released in 2009; current version as of May, 2019 is 4.0.8
- Serves and receives Json documents but stores them as BSON (Binary Json)
- MongoDB is CP with a flexible Consistency model.
- High profile users include Baidu, Foresquare, Craigslist, SourceForge and more.
- Powerful support for Geospatial use cases.





- ☐ Database
- ☐ Collection
- ☐ Shards
- ☐ Replication
- ☐ Document
- ☐ Field
- ☐ Data types
- ☐ GridFS

Movielens





Indexes

A thin vertical black line is positioned to the left of the text.

Document-oriented Database Wrap-up